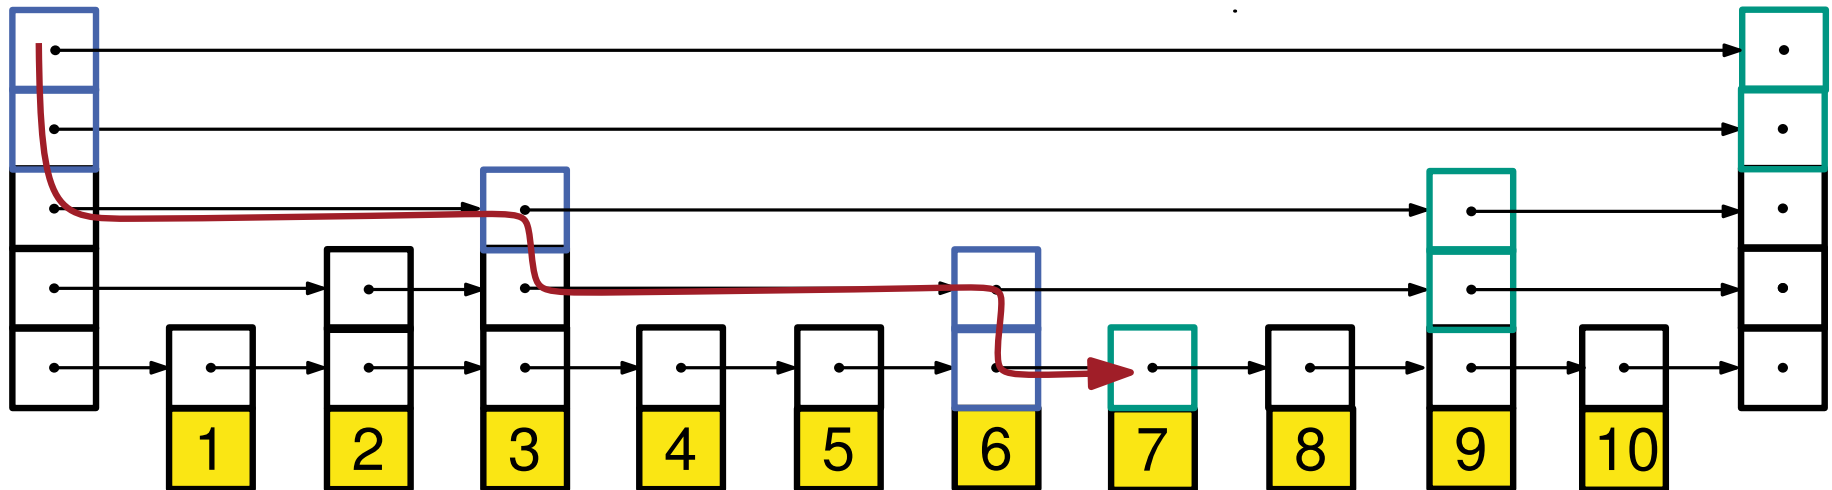


# Concurrent Skiplists: A Probabilistic Alternative to Balanced Trees

A Practical Course · Efficient Parallel C++ · March 3, 2022  
Manuel Haag


INSTITUTE OF THEORETICAL INFORMATICS · ALGORITHMICS GROUP



# What is a Skiplist?

**Goal:** dynamic ordered set

**Operations:** search, insert, remove




# What is a Skiplist?

**Goal:** dynamic ordered set

**Operations:** search, insert, remove

- bottom layer is linked list
- create nodes with random heights  $\rightarrow$  “fast lane” allows skipping nodes
- $P(\text{increase height by } k) = p^k(1 - p) \sim \text{Geo}(1 - p)$




# What is a Skiplist?

**Goal:** dynamic ordered set

**Operations:** search, insert, remove

- bottom layer is linked list
- create nodes with random heights  $\rightarrow$  “fast lane” allows skipping nodes
- $P(\text{increase height by } k) = p^k(1 - p) \sim \text{Geo}(1 - p)$



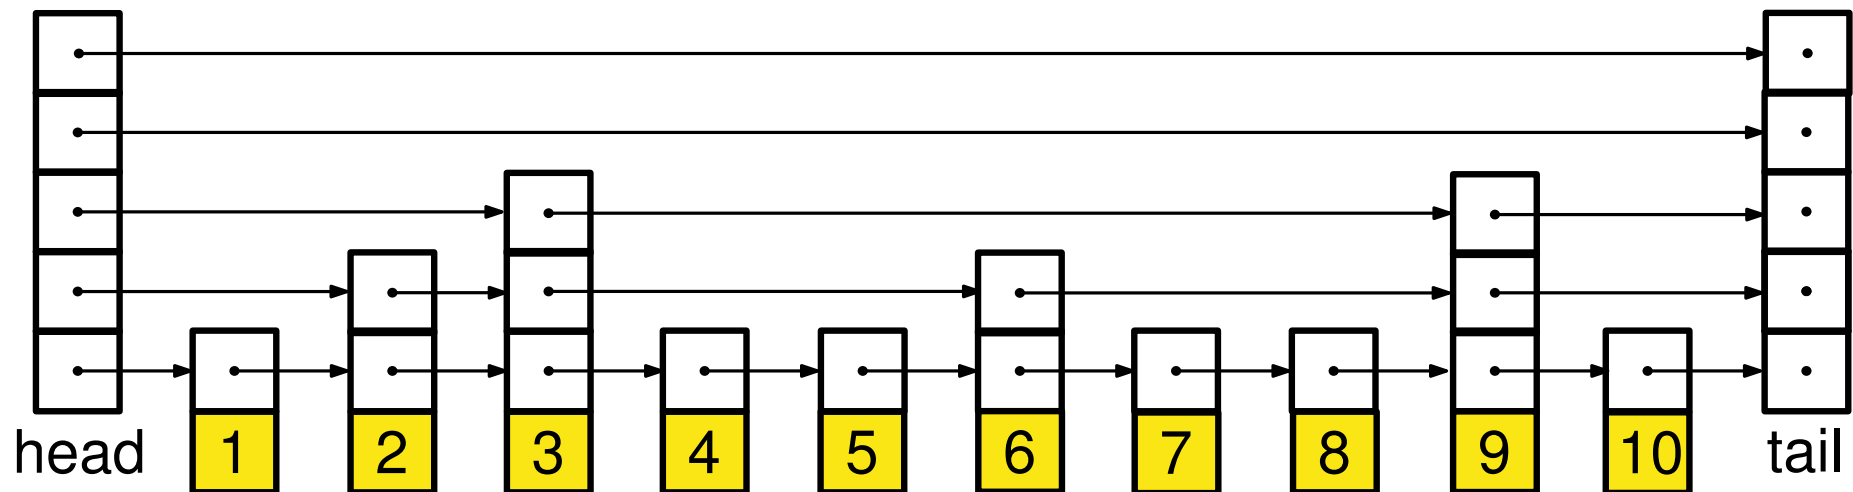
**Complexity:**  $\mathbb{E}[\text{cost}] = \frac{1}{p} \log_{\frac{1}{p}}(n) \in \mathcal{O}(\log n)$

**Storage:**  $\mathbb{E}[\text{height}] = \frac{1}{1-p} \in \mathcal{O}(n)$

# Extension to Indexable Skiplist

**Goal:** dynamic ordered set

**Operations:** search, insert, remove , **index**, **rank**

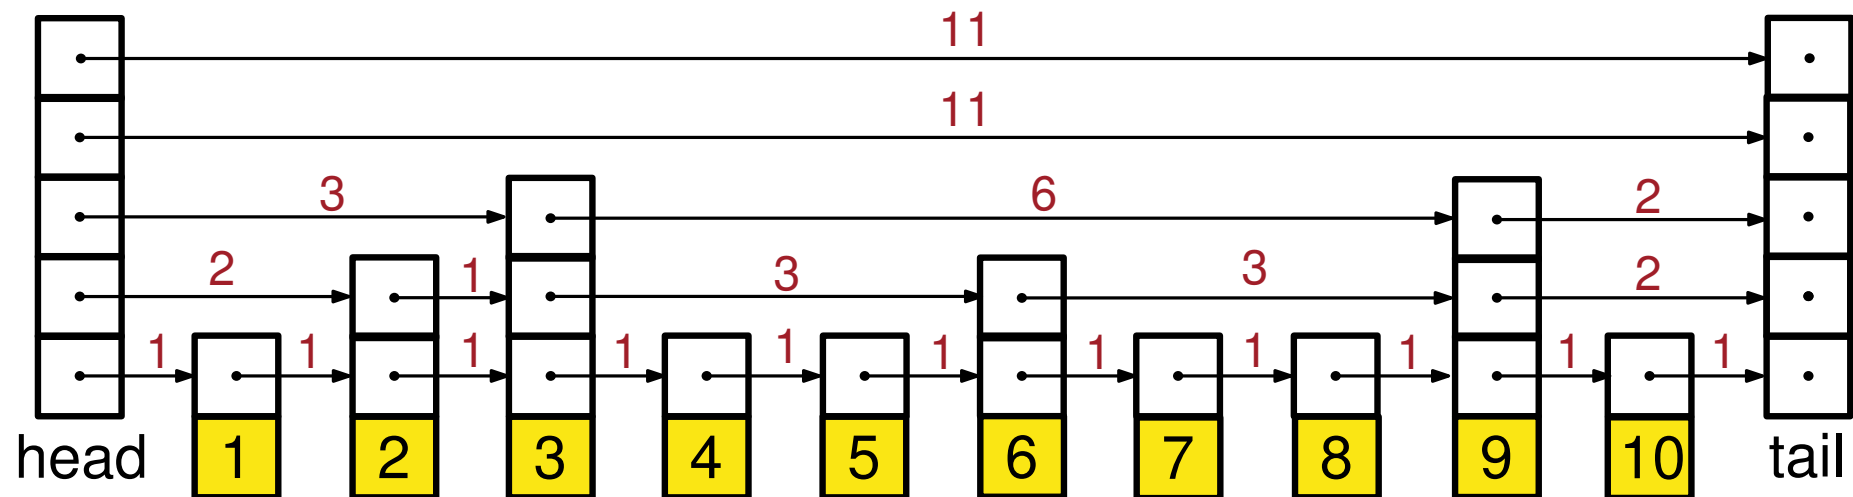


# Extension to Indexable Skiplist

**Goal:** dynamic ordered set

**Operations:** search, insert, remove, **index**, **rank**

- store length of pointers



# Motivation

- up to 5 supported operations
- probabilistic balancing
- easier to parallize than balanced trees

# Motivation


- up to 5 supported operations
- probabilistic balancing
- easier to parallize than balanced trees

## Exercise: Threadsafe Skiplist

- lockbased
- lockfree
- lockbased/lockfree + indexable
- evaluation




# Central Operation: find




# Central Operation: find

- find: traverse skiplist and document preds and succs (  $\neq$  search)
- basis for other operations



# Central Operation: find

- find: traverse skiplist and document preds and succs (  $\neq$  search)
- basis for other operations



- no locks are acquired  $\rightarrow$  no contention
- leads to inconsistencies  $\rightarrow$  validate and retry

# Implementation Highlights: Lock-Variant

- validate: preds and succs consistent and are not removed
- flags: being deleted, fully linked
- use of Spinlocks (faster than Mutex in preliminary experiments)

# Implementation Highlights: Lock-Variant

- validate: preds and succs consistent and are not removed
- flags: being deleted, fully linked
- use of Spinlocks (faster than Mutex in preliminary experiments)

```
auto try_insert_at = [&](int j) {  
    if(validate(j)) {  
        preds[j] -> lock.lock();  
        if(validate(j)) {  
            new_node -> next[j] = succs[j];  
            preds[j] -> next[j] = new_node;  
            preds[j] -> lock.unlock();  
            return true;  
        }  
        preds[j] -> lock.unlock();  
    }  
    find(preds, succs, insert_key); //get new preds and succs  
    return false;  
};
```

# Implementation Highlights: Lock-Variant


- validate: preds and succs consistent and are not removed
- flags: being deleted, fully linked
- use of Spinlocks (faster than Mutex in preliminary experiments)

```
auto try_insert_at = [&](int j) {  
    if(validate(j)) {  
        preds[j] -> lock.lock();  
        if(validate(j)) {  
            new_node -> next[j] = succs[j];  
            preds[j] -> next[j] = new_node;  
            preds[j] -> lock.unlock();  
            return true;  
        }  
        preds[j] -> lock.unlock();  
    }  
    find(preds, succs, insert_key); //get new preds and succs  
    return false;  
};
```

- thread to succeed at first level links remaining levels
- new\_node accessible after first linking → lock until the end
- remove similar


# Implementation Highlights: Lockfree-Variant

**Problem:** no validation and can only change one pointer atomically



# Implementation Highlights: Lockfree-Variant


**Problem:** no validation and can only change one pointer atomically






# Implementation Highlights: Lockfree-Variant

**Problem:** no validation and can only change one pointer atomically




3. removed




2. removed, but 3. inserted again!

# Implementation Highlights: Lockfree-Variant

**Problem:** no validation and can only change one pointer atomically



3. removed




2. removed, but 3. inserted again!


**Solution:** only remove from left to right

# Implementation Highlights: Lockfree-Variant

**Problem:** no validation and can only change one pointer atomically



3. removed



2. removed, but 3. inserted again!

**Solution:** only remove from left to right

- mark outgoing pointers of removed nodes
- unlink nodes lazy in `find`

# Implementation Highlights: Lockfree-Variant

**Trick:** last 16 bits of 64bit-pointer are not used

→ store bool in pointer address

→ allows atomically setting pointer and marked flag

# Implementation Highlights: Lockfree-Variant

**Trick:** last 16 bits of 64bit-pointer are not used

→ store bool in pointer address

→ allows atomically setting pointer and marked flag

```
template<class T>
class MarkableReference
{
private:
    uintptr_t val;
    static const uintptr_t mask = 1;
public:
    MarkableReference(T* ref = NULL, bool mark = false) {
        val = ((uintptr_t)ref & ~mask) | (mark ? 1 : 0);
    }
    T* getRef() const { return (T*)(val & ~mask); }
    bool getMark() const { return (val & mask); }
    T *operator->() const { return (T*)(val & ~mask); }
};
```

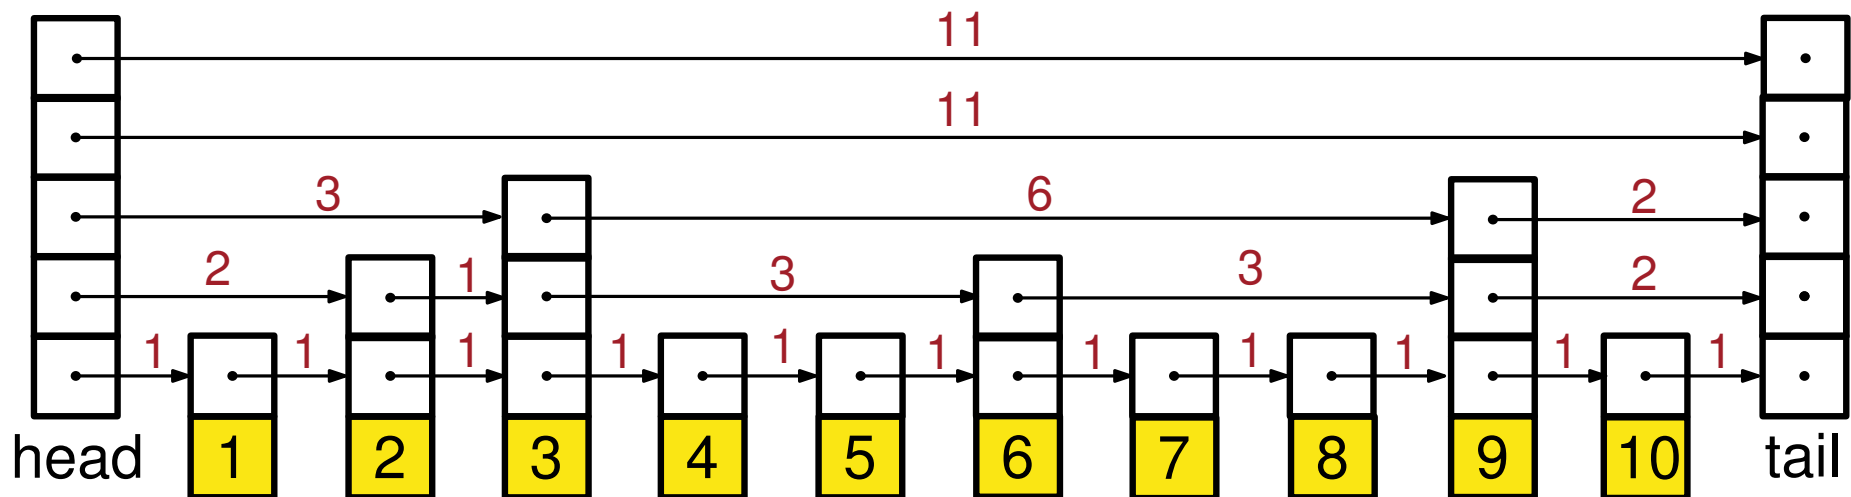
# Implementation Highlights: Lockfree-Variant

unlink in find:

```
// pred -> cur -> succ
while(succ.getMark()) {
    MarkPtr expect = {cur.getRef(), false};
    snip = (pred -> next[i].compare_exchange_strong(expect, {succ.getRef(), false}));
    if(!snip) {goto retry;} //restart find
    cur = pred -> next[i];
    succ = cur -> next[i];
}
```

# Implementation Highlights: Indexable-Variant

- hard to maintain lengths, since all length above change
- suboptimal solution: recompute lengths before rank and index operation



# Garbage Collection

- `shared_ptr` allows dynamic deallocation  
→ 10 times slower due to counting references




# Garbage Collection

- `shared_ptr` allows dynamic deallocation
  - 10 times slower due to counting references
- simple solution: store removed nodes in garbage queues and clean up at the end
  - memory will go out for long time usage
  - ok for short time usage


# Evaluation

Benchmark on 6 core laptop: 5 runs, 20% insert, 20% remove, 60% search, Key = int




# Evaluation

Benchmark on 6 core laptop: 5 runs, 20% insert, 20% remove, 60% search, Key = int



# Evaluation

Benchmark on 6 core laptop: 5 runs, 20% insert, 20% remove, 60% search, Key = int



- [1] W. Pugh. Skip lists: a probabilistic alternative to balanced trees. *ACM Transactions on Database Systems*, 33(6):668–676, 1990.
- [2] Y. Lev, M. Herlihy, V. Luchangco, and N. Shavit. A Simple Optimistic Skiplist Algorithm. *Fourteenth Colloquium on structural information and communication complexity (SIROCCO) 2007* pp. 124–138, June 5–8, 2007, Castiglioncello (LI), Italy.
- [3] Herlihy, Y. Lev, and N. Shavit. A lock-free concurrent skiplist with wait-free search. *Unpublished Manuscript*, Sun Microsystems Laboratories, Burlington, Massachusetts, 2007.
- [4] *The Art of Multiprocessor Programming*, Maurice Herlihy, Nir Shavit.

# Summary

## Lockbased

- lockfree find + validation
- lock levels individually
- first successful thread inserts/removes node completely

## Lockfree

- markable pointer
- delay remove to find

## Indexable

- recompute lengths before rank and index

## Garbage Collection

- shared\_ptr slow, thus deallocation of all deleted nodes at the end